

NoSQL Databases

Trainer: Mr. Nilesh Ghule

Agenda

- Introduction
- RDBMS vs NoSQL
- Scaling
- CAP theorem
- Advantages
- Limitations
- Applications
- Key-value database
- Column-oriented database
- Graph database
- Document database



Database

- A database is an organized collection of data, generally stored and accessed electronically from a computer system.
- A database refers to a set of related data and the way it is organized.
- Database Management System
 - Software that allows users to interact with one or more databases and provides access to all of the data contained in the database

- Types

- ✓ RDBMS

language: SQL

Oracle, MS-SQL, MySQL, DB2, PostgreSQL, etc.

Tables (rows + columns)

- ✓ NoSQL

language: no fixed lang, no fixed syntax - db specific.
Mongo, Cassandra, HBase, Redis, DynamoDB, Neo4J, etc.

- ✓ NewSQL

Structure: db specific.

high performance, scalability

high performance + SQL

(like NoSQL)

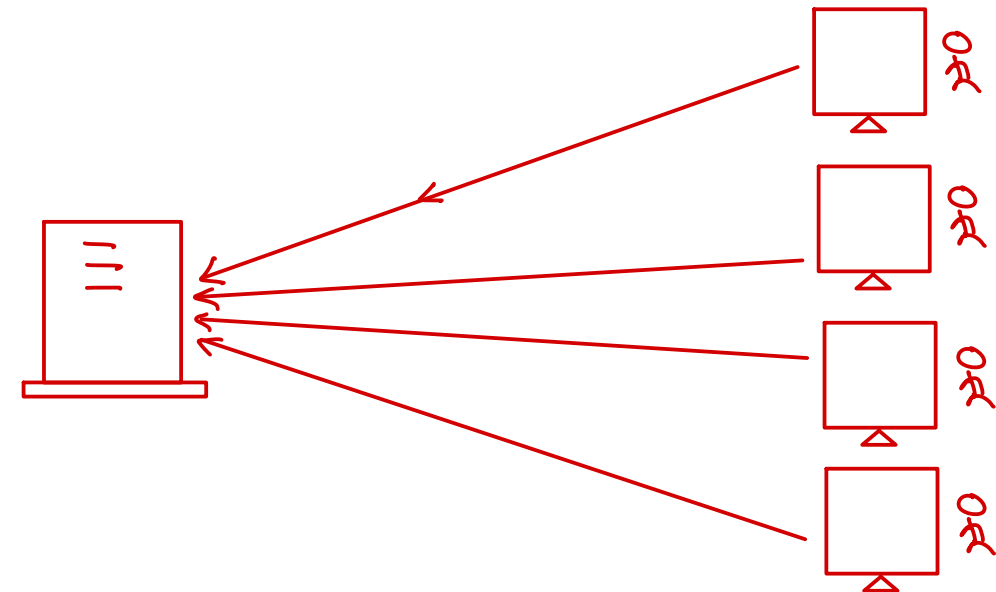
(like RDBMS)

HarperDB, WHDB, memSQL, ..



RDBMS

- The idea of RDBMS was borne in 1970 by E. F. Codd. *→ Codd's rules: 1+12*
- Structured and organized data *→ Fixed Schema → Create table: column names
+ Column types
+ Constraints*
- Structured query language (SQL)
- DML, DQL, DDL, DTL, DCL.
- Data and its relationships are stored in separate tables. *→ system tables*
- Tight Consistency
- Based on Codd's rules
- ACID transactions.
 - Atomic
 - Consistent
 - Isolated
 - Durable



NoSQL

- Refer to non-relational databases
- Stands for Not Only SQL
- Term NoSQL was first used by Carlo Strozzi in 1998.
- No declarative query language
- No predefined schema, Unstructured and unpredictable data
- Eventual consistency rather ACID property
- Based on CAP Theorem Consistency, Availability & Partition Tolerance.
- Prioritizes high performance, high availability and scalability
- BASE Transaction
 - Basically Available → service running 24x7
 - Soft state → data is arranged internally without any external inputs.
 - Eventual consistency

Structured
Fixed - no & cols

Schema : Flexible, Evolve.
(semi-structured)
XML JSON

Unstructured
media
✓ images
✓ audio
✓ video
text

1970-1990: RDBMS
198x: Internet
1990: WWW
1995: Applets in Java

Anti-SQL
1998: NoSQL
Carlo

2004: Last.fm
↓
International Conferences
#nosql

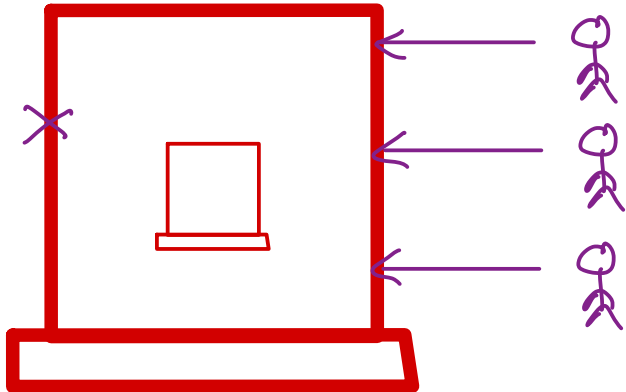
Not Only SQL



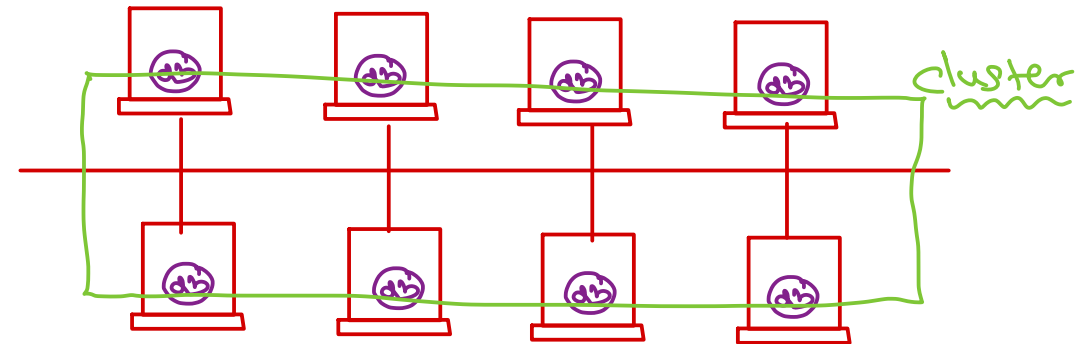
Scaling

- Scalability is the ability of a system to expand to meet your business needs.
- E.g. scaling a web app is to allow more people to use your application.
- Types of scaling
 - Vertical scaling: Add resources within the same logical unit to increase capacity. E.g. add CPUs to an existing server, increase memory in the system or expanding storage by adding hard drives.
 - Horizontal scaling: Add more nodes to a system. E.g. adding a new computer to a distributed software application. Based on principle of distributed computing.
- NoSQL databases are designed for Horizontal scaling. So they are reliable, fault tolerant, better performance (at lower cost), speed.

users ↑
data ↑
time ↓

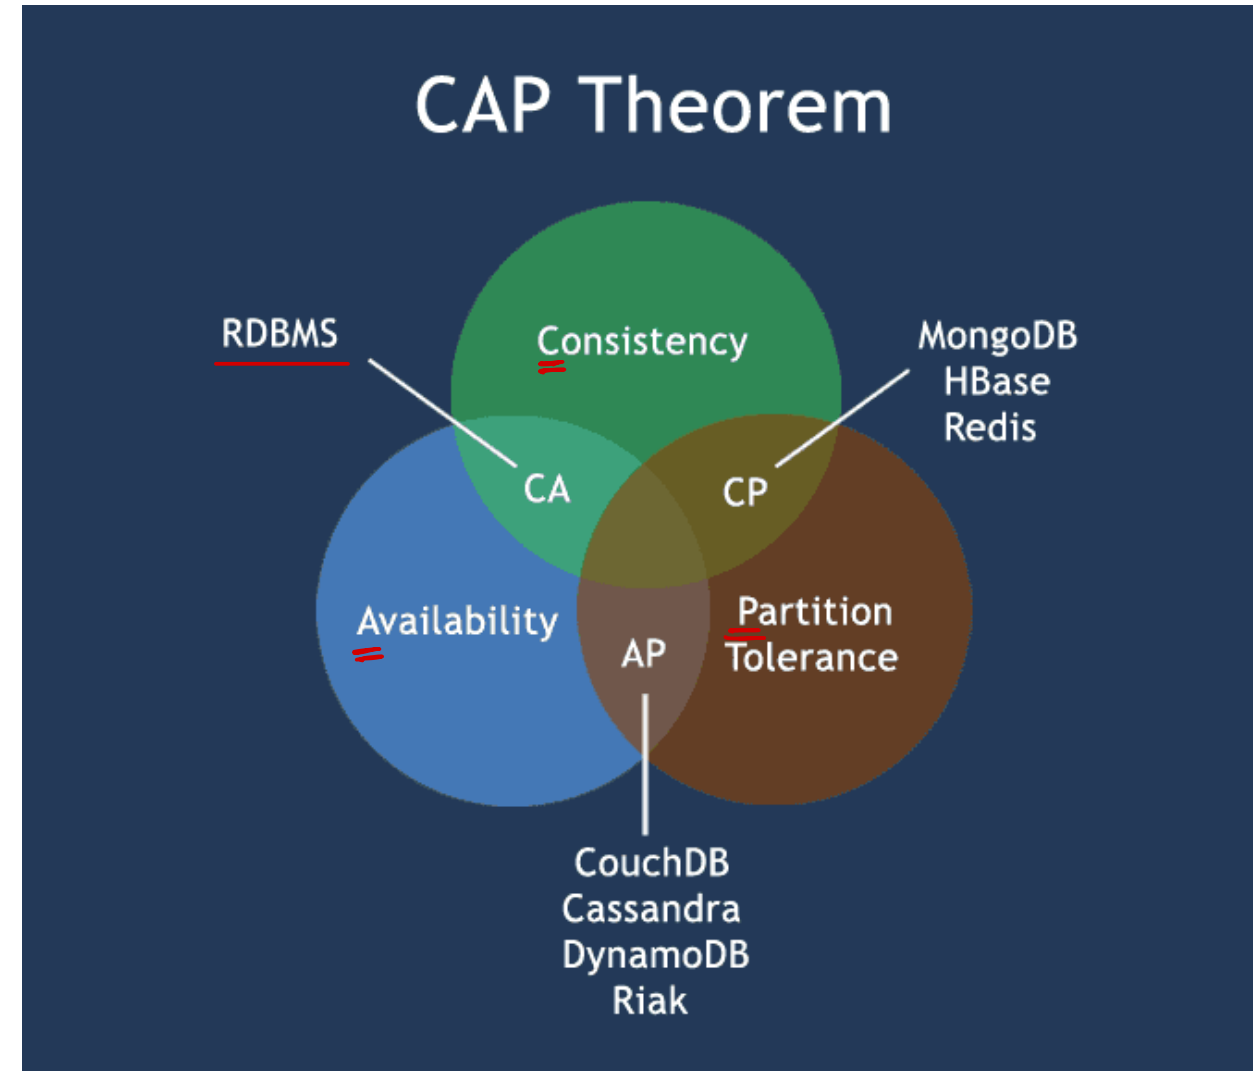


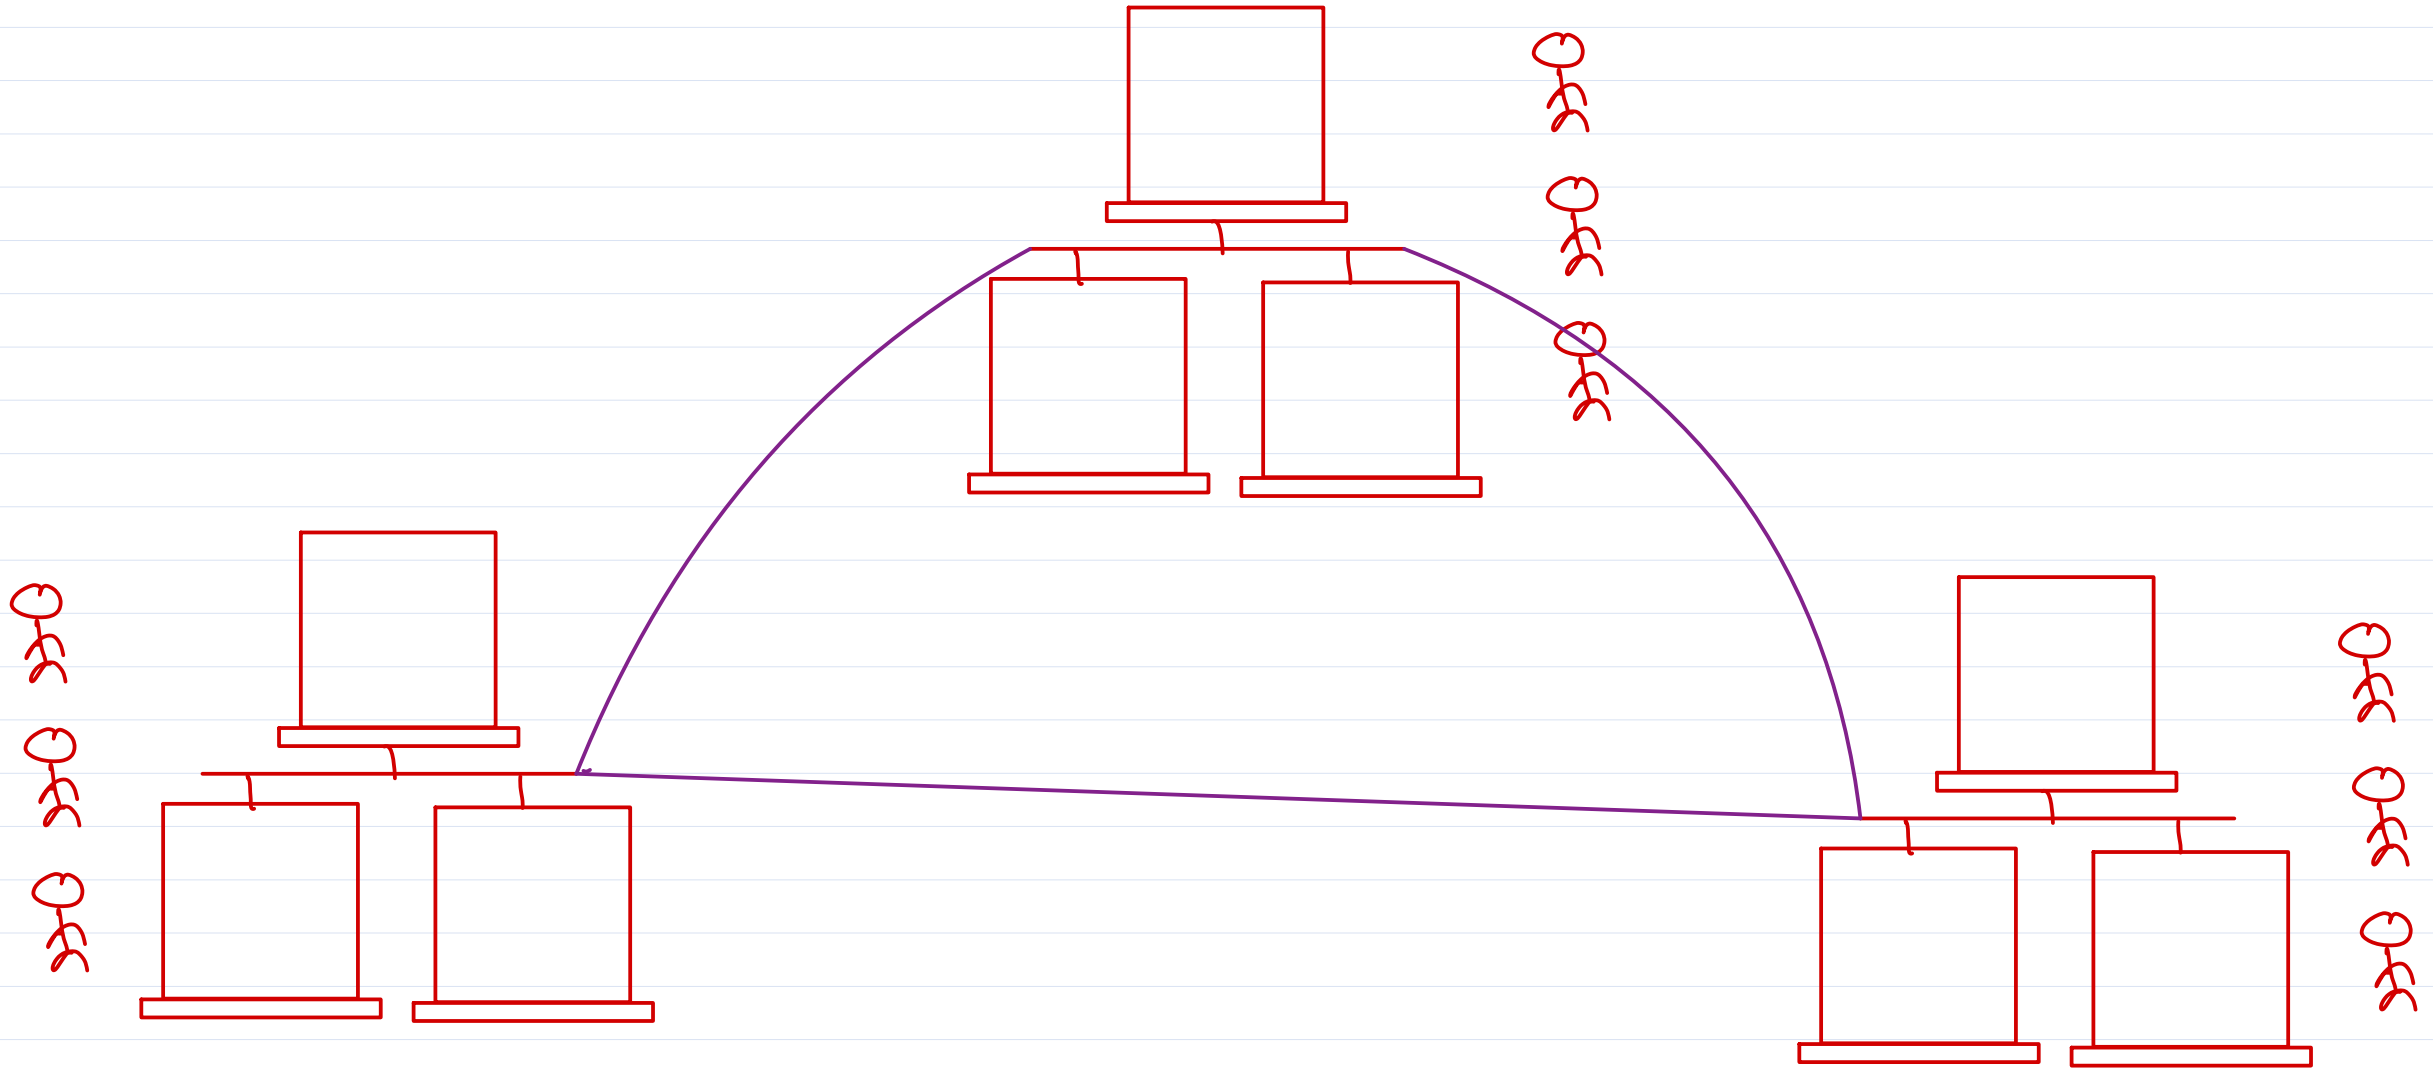
Availability ↑



CAP (Brewer's) Theorem

- **Consistency** - Data is consistent after operation. After an update operation, all clients see the same data.
- **Availability** - System is always on (i.e. service guarantee), no downtime.
- **Partition Tolerance** - System continues to function even the communication among the servers is unreliable.
- **Brewer's Theorem**
 - It is impossible for a distributed data store to simultaneously provide more than two out of the above three guarantees.





Advantages of NoSQL

- High scalability
 - This scaling up approach fails when the transaction rates and fast response requirements increase. In contrast to this, the new generation of NoSQL databases is designed to scale out (i.e. to expand horizontally using low-end commodity servers).
- Manageability and administration
 - NoSQL databases are designed to mostly work with automated repairs, distributed data, and simpler data models, leading to low manageability and administration.
- Low cost
 - NoSQL databases are typically designed to work with a cluster of cheap commodity servers, enabling the users to store and process more data at a low cost.
- Flexible data models
 - NoSQL databases have a very flexible data model, enabling them to work with any type of data; they don't comply with the rigid RDBMS data models. As a result, any application changes that involve updating the database schema can be easily implemented.



Disadvantages of NoSQL

- **Maturity**
 - Most NoSQL databases are pre-production versions with key features that are still to be implemented. Thus, when deciding on a NoSQL database, you should analyse the product properly to ensure the features are fully implemented and not still on the To-do list.
- **Support**
 - Support is one limitation that you need to consider. Most NoSQL databases are from start-ups which were open sourced. As a result, support is very minimal as compared to the enterprise software companies and may not have global reach or support resources.
- **Limited Query Capabilities**
 - Since NoSQL databases are generally developed to meet the scaling requirement of the web-scale applications, they provide limited querying capabilities. A simple querying requirement may involve significant programming expertise.
- **Administration**
 - Although NoSQL is designed to provide a no-admin solution, it still requires skill and effort for installing and maintaining the solution.
- **Expertise**
 - Since NoSQL is an evolving area, expertise on the technology is limited in the developer and administrator community.



Applications

- When to use NoSQL?

- Large amount of data (TBs)
- Many Read/Write ops
- Economical Scaling
- Flexible schema

- Examples:

- ✓ Social media
- ✓ Recordings
- ✓ Geospatial analysis
- ✓ Information processing

- When Not to use NoSQL?

- Need ACID transactions ✗
- Fixed multiple relations ✗
- Need joins ✗
- Need high consistency ✗

- Examples

- Financial transactions ✗
- Business operations ✗



RDBMS vs NoSQL

	RDBMS	NoSQL
Types	All types support SQL standard	Multiple types exists, such as document stores, key value stores, column databases, etc
History	Developed in 1970	Developed in 2000s
Examples	SQL Server, Oracle, MySQL	MongoDB, HBase, Cassandra, Redis, Neo4J
Data Storage Model	Data is stored in rows and columns in a table, where each column is of a specific type	The data model depends on the database type. It could be Key-value pairs, documents etc
Schemas	Fixed structure and schema	Dynamic schema. Structures can be accommodated
Scalability	Scale up approach is used	Scale out approach is used
Transactions	Supports ACID and transactions	Supports partitioning and availability
Consistency	Strong consistency	Dependent on the product [Eventual Consistency]
Support	High level of enterprise support	Open source model
Maturity	Have been around for a long time	Some of them are mature; others are evolving



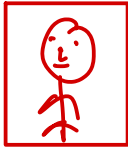
NoSQL database

- NoSQL databases are non-relational.
- There is no standardization/rules of how NoSQL database to be designed.
- All available NoSQL databases can be broadly categorized as follows:
 - ✓ • Key-value databases
 - ✓ • Column-oriented databases
 - ✓ • Graph databases
 - ✓ • Document oriented databases



Key-value database

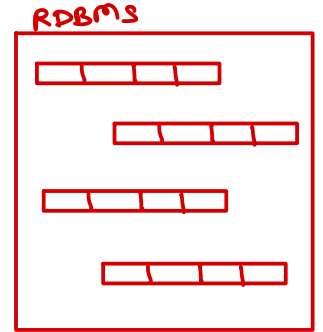
- Based on Amazon's Dynamo database.
- For handling huge data of any type.
- Keys are unique and values can be of any type i.e. JSON, BLOB, etc.
- Implemented as big distributed hash-table for fast searching.
- Example: redis, dynamodb, riak, ... *Oracle KV Store, etc.*

key value
dbda/1/name = James Bond
dbda/1/addr = Universal Exports, London
dbda/1/photo = 
dbda/1/age = 30



Column-oriented databases

- Values of columns are stored contiguously.
- Better performance while accessing few columns and aggregations.
- Good for data-warehousing, business intelligence, CRM, ...
- Examples: hbase, cassandra, bigtable, ...



HBase : books : id, name, author, subject, price,

	name	author	subject	price
1	let us c	abc	c	200.0
2	java lang	pqr	java	300.0
3	python lang	xyz	python	250.0
	~	~	~	~
	~	~	~	~
	~	~	~	~

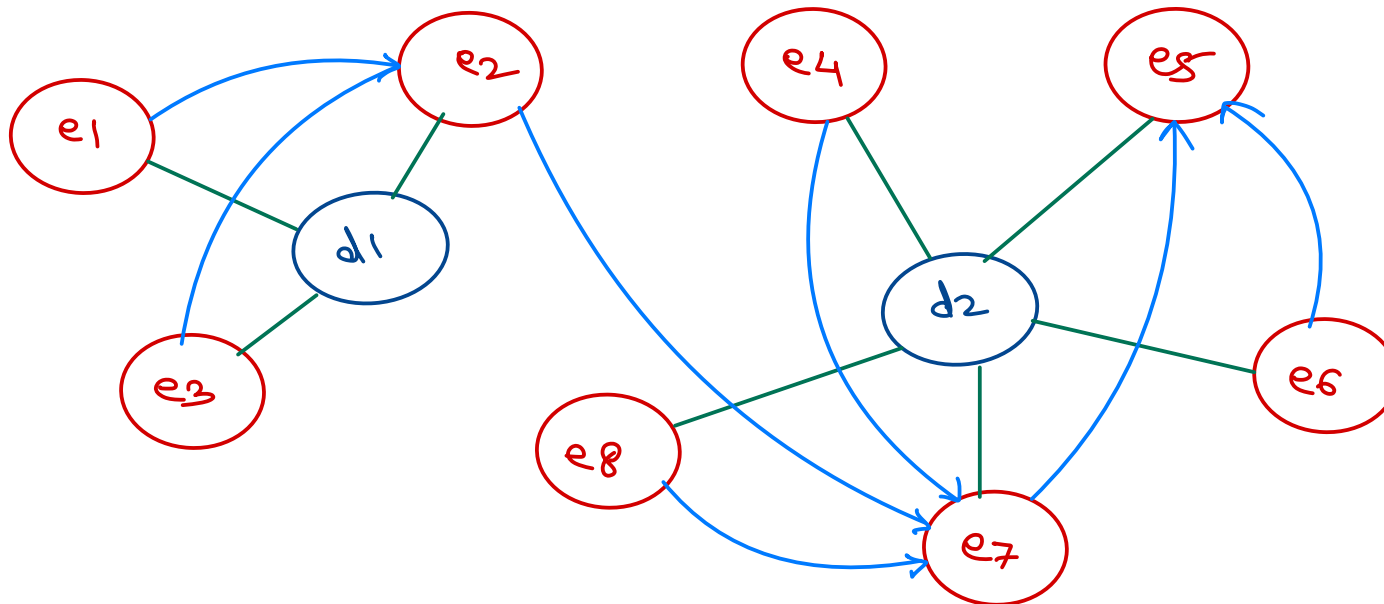
Select subject, sum (price) from
books group by subject;



Graph databases

← (data structure)

- Graph is collection of vertices and edges (lines connecting vertices).
- Vertices keep data, while edges represent relationships.
- Each node knows its adjacent nodes. Very good performance, when want to access all relations of an entity (irrespective of size of data).
- Examples: Neo4J, Titan, ...



Document oriented databases

- Document contains data as key-value pair as JSON or XML.
- Document schema is flexible & are added in collection for processing.
- RDBMS tables → Collections
- RDBMS rows → Documents
- RDBMS columns → Key-value pairs in document
- Examples: MongoDB, CouchDb, ...





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

