



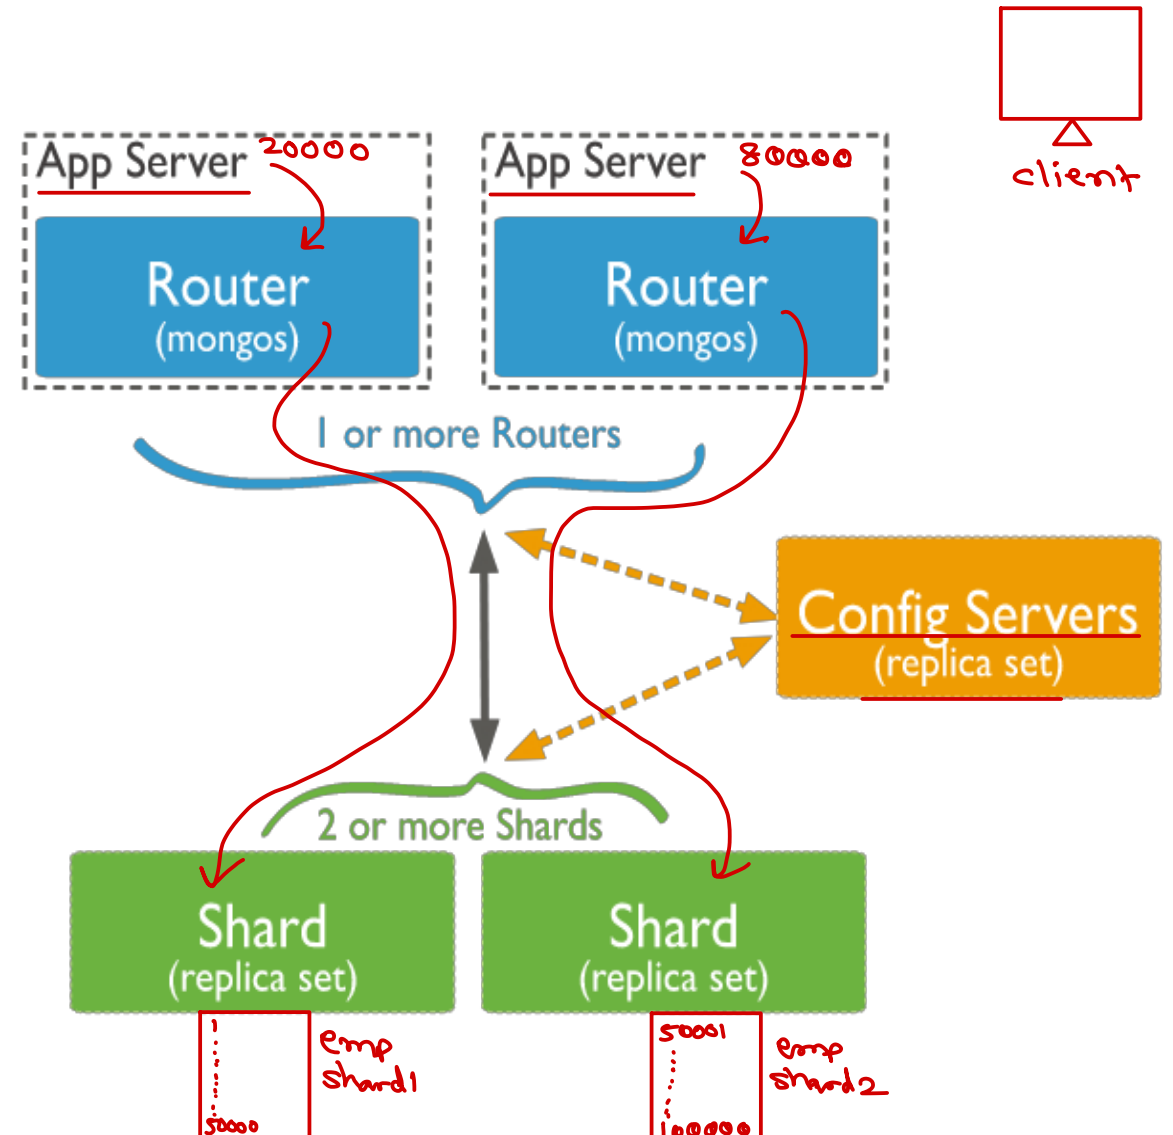

MongoDb Databases

Trainer: Mr. Nilesh Ghule



Mongo - Sharding

- Sharding is a method for distributing large data across multiple machines.
- This is mongodb approach for horizontal scaling/scaling out.
- shard: part of collection on each server (replica set).
- mongos: query router between client & cluster.
- config servers: metadata & config settings of cluster.




Mongo - Sharding

- Collections can be sharded across the servers based on shard keys.
- Shard keys:
 - Consist of immutable field/fields that are present in each document
 - Only one shard key. To be chosen when sharding collection. Cannot change shard key later.
 - Collection must have index starting on shard key.
 - Choice of shard key affect the performance.
- Advantages:
 - Read/Write load sharing
 - High storage capacity
 - High availability



Mongo - Sharding

- Sharding strategies:
 - Hashed sharding
 - MongoDB compute hash of shard key field's value.
 - Each chunk is assigned a range of docs based on hashed key.
 - Even data distribution across the shards. However range-based queries will target multiple shards.
 - Ranged sharding
 - Divides data into ranges based on shard key values.
 - mongos can target only those shards on which queried range is available.
 - Efficiency of sharding is based on choosing proper shard key.





Sunbeam Infotech

Redis



Redis - Introduction

→ key-value

- REmote Dictionary Server
- In-memory persistent open-source key-value store developed in 2009.
- Redis is maintained and developed by Salvatore Sanfilippo.
- Based on data structures: strings, hashes, sets, lists, sorted sets, geospatial indexes, hyperloglogs. ✓ ✓ ✓ ✓ ✓
- Application/Uses:
 - Advanced key/value store as NoSQL.
 - Used as memory cache to improve application performance.
 - Message broker for real time message notifications.
 - Easy and efficient implementation of Data structures.

Redis - Features



- Speed: 110000 SET/s and 81000 GET/s on entry-level Linux system. ↙
- Pipeline: Multiple commands execution for faster execution.
- Persistence: Whole data accessed from memory, asynchronously persisted on disk with flexible policies.
- Data Structure: Based on data structures like Strings, Hashes, Sets, ...
- Atomic operations: Data is manipulated atomically by multiple clients.
- Supported Languages: Drivers available for C/C++, Java, Python, R, PHP, ...
- Master/Slave replication: Easy config and fast execution. ✓ ✓ ✓ ✓
- Sharding: Distributing across cluster. Based on client driver capability.
- Portable: Developed in C. Work on all UNIX variants. Not supported on Win.

Redis - Highlights

- Key-value DB, where values can store complex data types with atomic ops.
- Value types are basic data structures made available to programmers without layers of abstraction.
- It is in-memory but persistent store i.e. whole database is maintained in server RAM, only changes are updated on disk for backup.
- The data storage in disk is in append-only data files.
- Maximum data size is limited to the RAM size.
- On modern systems if Redis is going out of memory, it will start swapping and slow down the system.
- Max memory limit can be configured to raise error on write or evict keys.

Redis - Installation

- Install: `sudo apt-get install redis-server redis-tools` ✓
- Run server: `sudo systemctl start redis` ✓
- Run client: `redis-cli` ✓
- `redis> ping` → `PONG`
- `redis> INFO`
- `redis> CONFIG GET *`
- `redis> CONFIG GET loglevel` ← *key*
- `redis> CONFIG SET loglevel notice`
 - loglevels: 0. debug, 1. verbose, 2. notice, 3. warning
- `redis> KEYS *`



Redis - Data Types & Commands

▪ Keys

- Any binary sequence as key i.e. any string to any binary file.
- Max key size is 512 MB. Very large key size is not good.
- Set up convention for key e.g. users:1001:posts.november

▪ Data Types:

- ✓ String: Basic type. (SET, GET, DEL)
- ✓ List: Ordered collection. (LPUSH, RPUSH, LPOP, RPOP, LREM, LRANGE).
- ✓ Set: Ordered collection. Unique values. (SADD, SMEMBERS, SISMEMBER).
- ✓ Sorted Set: Sorted collection. Unique. Each value have score value (float) for sorting. (ZADD, ZRANGE).
- ✓ Hashes: Object with multiple fields. (HMSET, HGETALL, HMGET)

Redis - Publish/Subscribe

- PSUBSCRIBE channel-pattern
 - receive notifications from given channels. e.g. b?g, b*g, b[ai]g
- PUBLISH channel "message"
 - send message to channel
- PUNSUBSCRIBE channel-pattern
 - stop receiving notifications from given channels.
- UNSUBSCRIBE channel
 - stop receiving notifications from given channel.
- PUBSUB *command*
 - monitor pub-sub subsystem
 - e.g. PUBSUB channels

Redis-Transactions & Pipeline

Transaction:

- Puts multiple commands in a queue and execute them at once.
- MULTI: begin transaction
- All commands after this are queued.
- EXEC: execute all commands from start of transaction
- DISCARD: discard all commands from start of transaction

Pipeline:

- Client sends multiple commands to server in a batch.
- Saves network round-trip each time.
- All commands may not execute in a transaction.
- echo -en "PING\r\nSET key value\r\nGET key\r\n" | nc localhost 6379

on server side

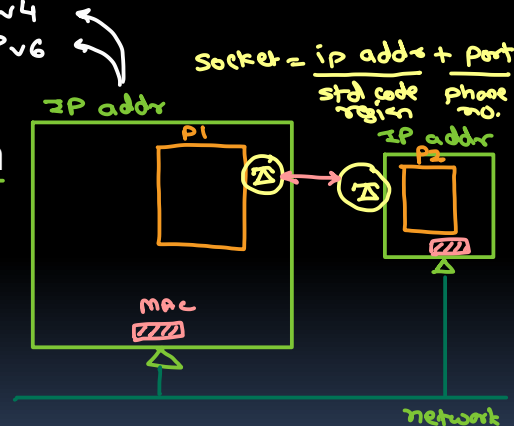
192.168.0.1

xxx.xxx.xxx.xxx

0-255

32-bit : IPv4

128-bit : IPv6




netcat

SunBeam Infotech



Oracle NoSQL - KVStore

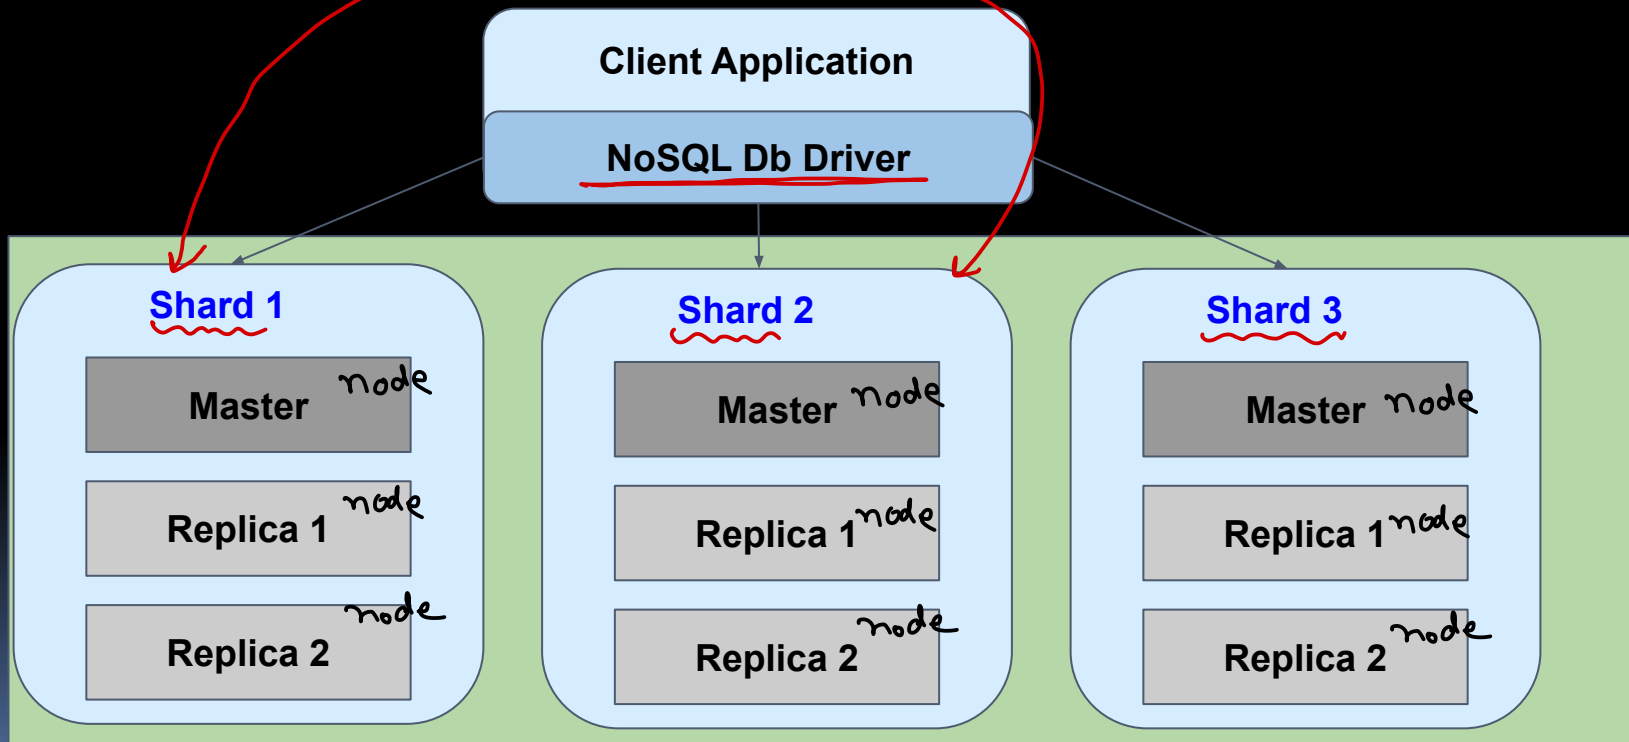


- Multi-terabyte distributed key-value pair storage → KV Store
- High performance, scalable, Eventual consistency, Durable.
- User defined read/write performance levels.
- Terminologies:
 - KV Pair → Key: Major & Minor keys, Value: byte array
 - KV Store → Container of KV pairs
 - Partition → Hashed Set of Records (on major keys)
 - Shard → Set of partitions. Group of machines for replication. Shard is chosen transparently i.e. auto selected by oracle nosql db.
 - Replication factor → Number of replicas. Default is 3.
 - Storage node → Physical machine for storing data (CPU+RAM+Disk).

Architecture

Key
/dbda/1/-/name
/dbda/1/-/addr
/dbda/2/-/name

value
James
UK
Bill



Consistency

- Related to update operation.
- Eventual consistency.
- Trade-off between : Speed & Availability
- Write transaction durability consists of Sync Policy & Replica Ack:
 - Sync Policy:
 - ✓ Sync (to disk) → Most Durable
 - ✗ Write No Sync (to OS buffer) → Moderate
 - ✗ No Sync (local log buffer - flush when convenient) → Fastest
 - Replica Ack Policy:
 - All → slower
 - Simple Majority (majority of nodes)
 - None → fastest

Consistency


- Read consistency:
 - Absolute (from Master)
 - Most Consistent : Most recent version
 - Time based (from replica within time-interval of Master)
 - Data of known version or later
 - Version (from replica with current/higher version of transaction token)
 - Recent data for given time
 - None (any replica)
 - Fastest : Can read stale data

Installation

- Download kv-ce-4.3.11.tar.gz and extract to some directory → kv-ce-4.3.11
- Edit ~/.bashrc
 - export KVHOME=<path to kv-ce-4.3.11>
 - export KVROOT=<path to kv-ce-4.3.11/kvroot>
- Start kvstore and test it.
 - java -jar \$KVHOME/lib/kvstore.jar kvlite -verbose -root \$KVROOT -store kvstore -host \$HOSTNAME -port 5000 -secure-config disable
 - java -jar \$KVHOME/lib/kvstore.jar ping -verbose -host \$HOSTNAME -port 5000
 - java -jar \$KVHOME/lib/kvstore.jar runadmin -verbose -host \$HOSTNAME -port 5000 -store kvstore

KV CLI :: kv ->

- show versions
 - show topology
 - verify
 - history
 - put kv -key <key> -value <value>
 - get kv -key <key>
 - get kv -key <key> -all
 - delete kv -key <key>
 - delete kv -key <key> -all
- Handwritten notes:*
- For put kv -key <key> -value <value>:
 - Handwritten: *major minor* above */dbda/1/-name*
 - Handwritten: *James* next to *-value <value>*
 - For get kv -key <key>:
 - Handwritten: */dbda/1/-name* next to the command
 - For get kv -key <key> -all:
 - Handwritten: */dbda/1/* next to the command



cassandra

Sunbeam Infotech



Introduction

- Google BigTable

- High performance data storage system built on GFS and other Google technologies.
- Master-slave architecture.
- One key, multiple values.
- Columnar, SSTable (Sorted String Table) Storage, Append-only, Memtable, Compaction.

Google File System (Distributed File System)
HDFS follow GFS concepts.


- Amazon DynamoDb

- Highly available and scalable key-value storage system.
- Decentralized peer to peer architecture.
- Compromise on consistency for better availability -- Eventual consistency.
- Consistent hashing, Gossip protocol, Replication, Read repair.

- Cassandra

- Inherited from BigTable and DynamoDb
- BigTable: Column families, Memtable, SSTable
- DynamoDb: Consistent hashing, Partitioning, Replication

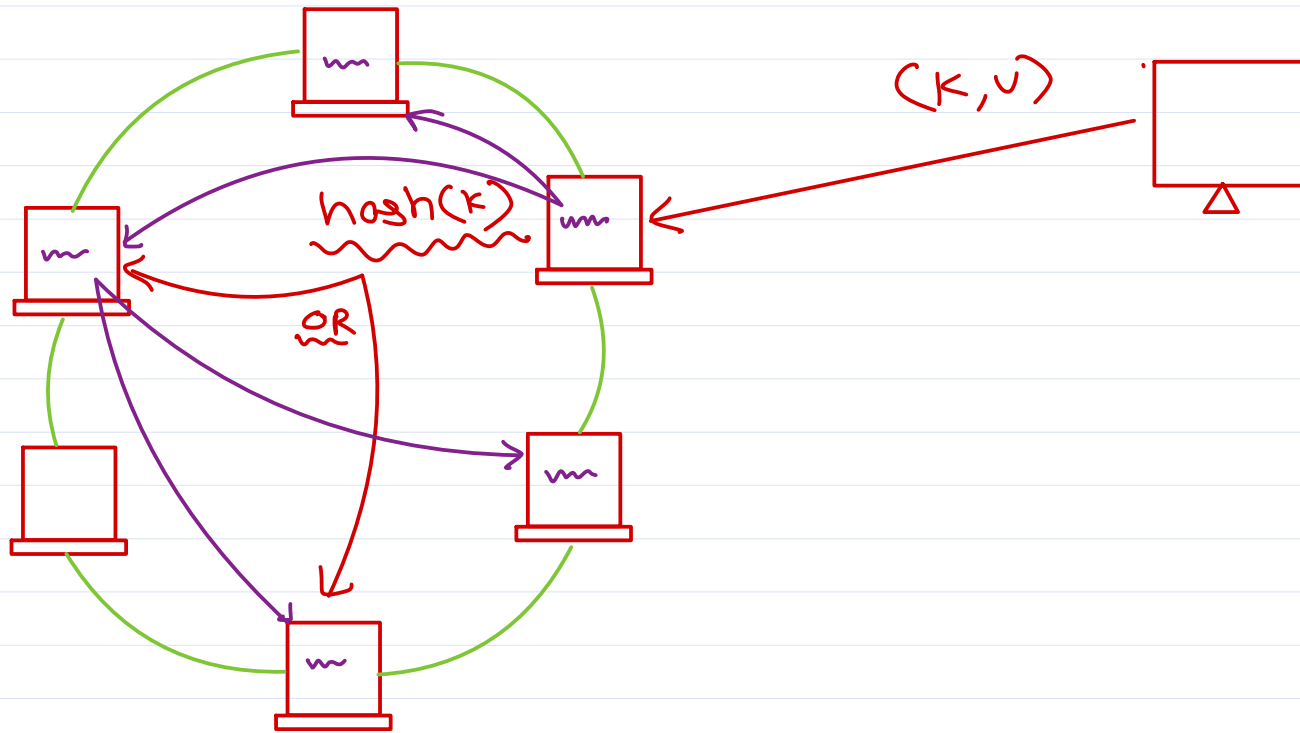




	ename	job	sal
1	ABC	1 CLERK	1 1000
2	POB	2 MGR	2 3000
3	XYZ	3 SALES	3 2500

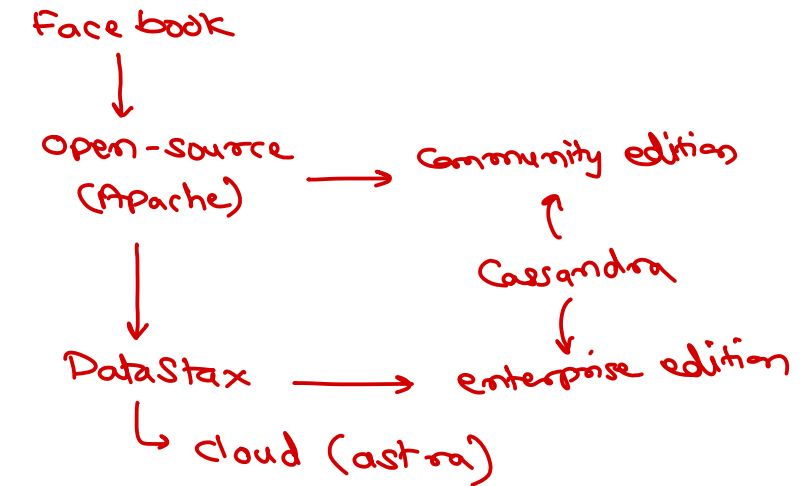
row id





Concept

- Developed by
 - Avinash Laxman (Co-inventor Amazon DynamoDb)
 - Prashant Malik (Technical Leader at Facebook).
- Goals:
 - Distributed NoSQL database (on commodity hardware)
 - Large amount of structured data
 - High availability
 - No single point of failure
- Basic data model is rows & columns.
- Column-oriented, Decentralized peer to peer & follow Eventual consistency.
- Datastax company develop and support commercial edition of Cassandra.



Cassandra Development

- Developed in Java.
- 2007-2008 - Developed at Facebook.
- July 2008 - Open sourced by Facebook.
- March 2009 - Apache Incubator project.
- February 2010 - Apache Top-level project.
- 2011 - version 0.8 - Added CQL.
- 2013 - version 2.0 - Added light-weight transactions, Triggers.
- 2015 - version 3.0 - Storage engine improved, Materialized views.
- 2020 - version 3.11 - Latest release.



Cassandra installation

- Prerequisite
 - Java 8 (Java 11 experimental)
- Can be installed through apt or yum tool (Ubuntu/CentOS).
- Manual installation
 - Download Cassandra 3.11.x (.tar.gz) and extract it.
 - set CASSANDRA_HOME to Cassandra directory.
 - set JAVA_HOME to JDK 8 directory.
 - Install python 2.7 (for cqlsh).
 - set CASSANDRA_HOME/bin into PATH variable
 - Start Cassandra
 - terminal 1> cassandra
 - terminal 2> cqlsh





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

