

cassandra

Sunbeam Infotech



Features

- Peer to peer architecture (no master-slave → no single point of failure).
- Linear scale performance Capacity \propto nodes
- High performance → high speed read/write
- Simplified deployment and maintenance → on linux
- Less expensive → horizontal scalability
- Supports multiple programming languages → Java, Python, ... & REST services
- Operational and Development simplicity
- Cloud availability → avail on all leading cloud vendors AWS, Azure, GCP, ...
- Ability to deploy across data-centers
- Fault tolerance
- Configurable consistency (tight or eventual) → tradeoff with speed
- Flexible data model → columnar storage. Number of columns can be added dynamically (super column)
- Column family store
- Schema-free



Limitations

- Aggregation operations not supported
- Range queries on partition key are not supported
- Not good for too many joins (Denormalized data preferred)
- Not suitable for transactional data - "C" poor from ACID → eventual consistency
- During compaction performance/throughput slow down
- Not designed for update-delete (is possible)



Performance

- Performance measures

- Throughput (Operations per second) → number of reads/writes per second (high)
- Latency → time required for single operation (low)

- Cassandra vs MySQL

- MySQL (more than 50 GB data)
 - Write speed: 300 ms
 - Read speed: 350 ms
- Cassandra (more than 50 GB data)
 - Write speed: 0.12 ms
 - Read speed: 15 ms



Applications

- Applications

- Product catalog/Playlist
- Recommendation/Personalization engine
- Sensor/IoT data
- Messaging/Time-series data
- Fraud detection

- Customers

- Facebook, Netflix, eBay, Apple, Walmart, GoDaddy

- Application requirements

- Store and handle time-series data
- Store and handle large volume of data
- Scale predictably (Linear Scaling)
- High availability

uuid - universally unique
identifier - 128 bit

timeuuid - uuid + timestamp
(helpful in sorting
on time field)



Architecture

- Commit Log

- Append only log of all mutations local to a node.
- Client data → commit log → memtable.
- Durability in the case of unexpected shutdown. On startup, any changes in log will be applied to tables.

- Memtable


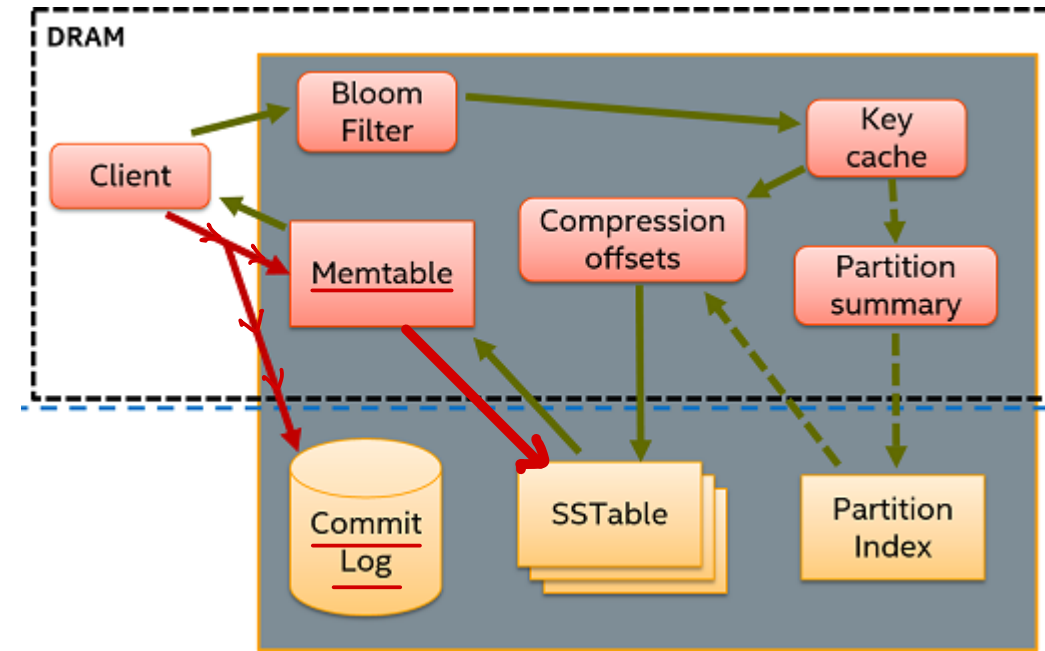
- In-memory structures to write Cassandra buffers.
- One active memtable per table.

- Sorted String Table

- Immutable data files for persisting data on disk.
- Multiple memtables merged into single SSTable.

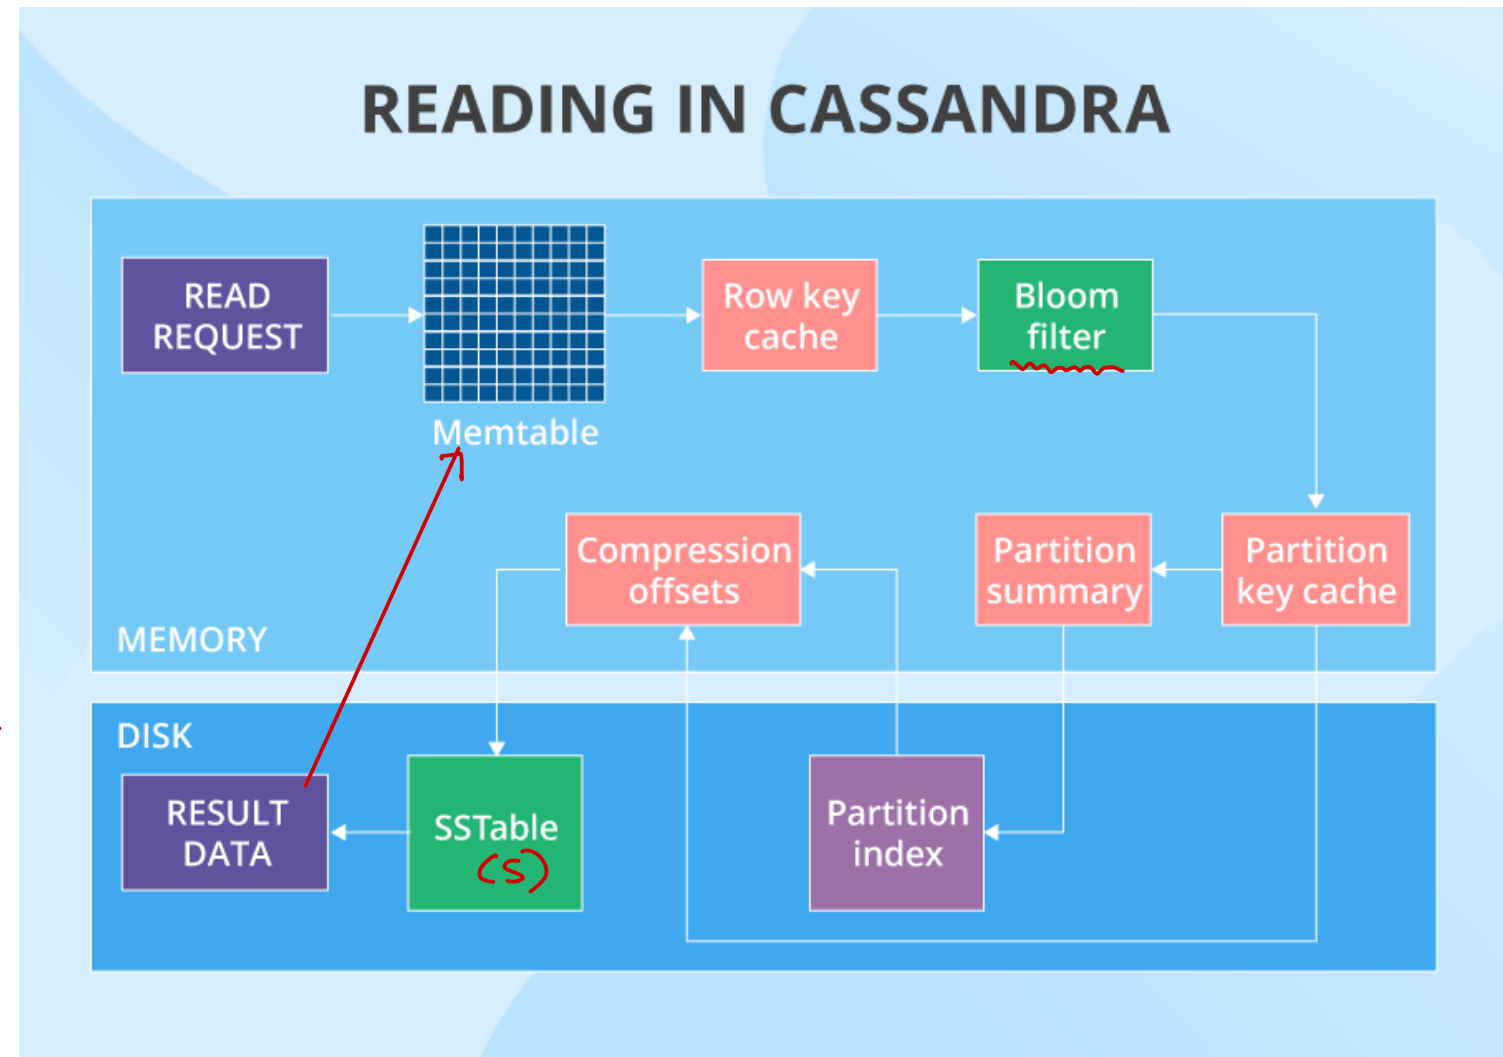
- LSM Tree (Log structure Merge Tree)

- Disk based data structure to provide low-cost indexing for a file, in which records are to be inserted at very high rate. also efficient in-mem structure.



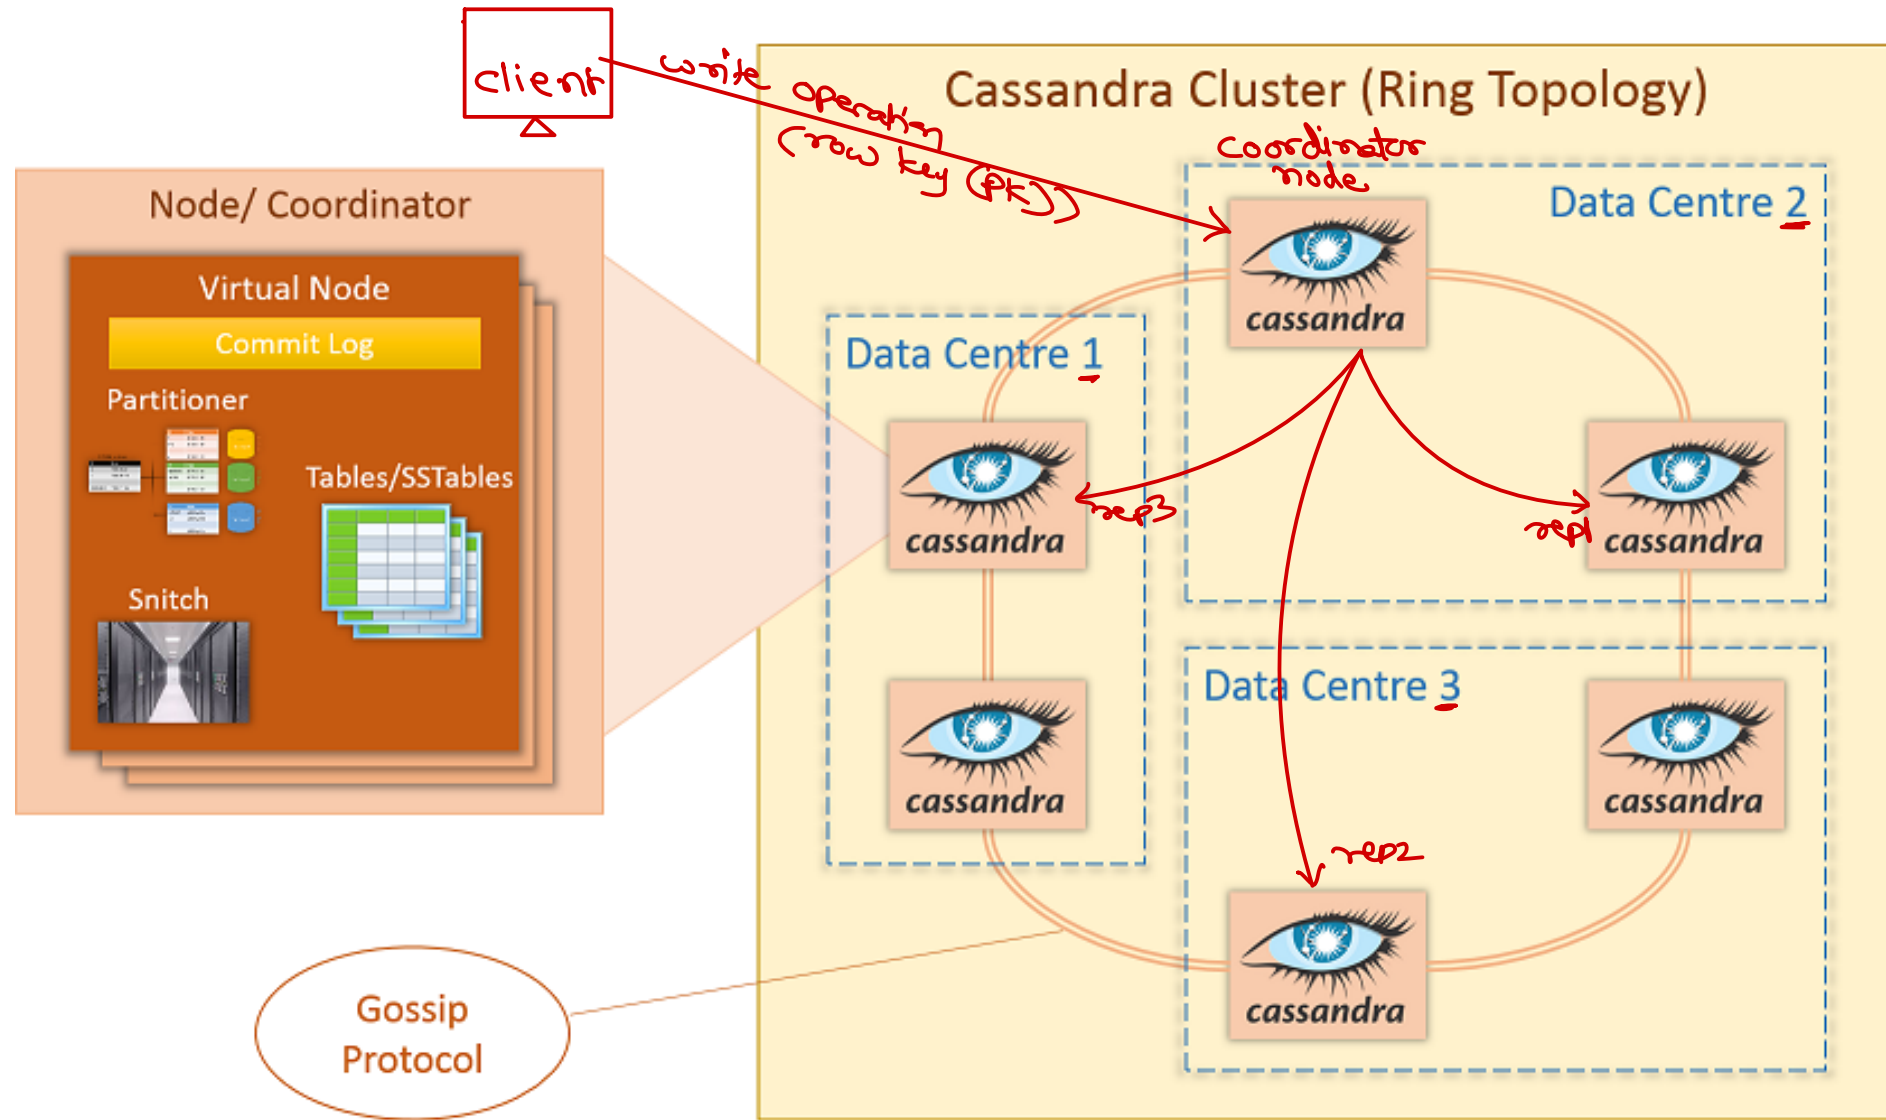
Architecture

- Bloom filter
 - In the read path, Cassandra merges data on disk (in SSTables) with data in RAM (in memtables).
 - To avoid checking every SSTable data file for the partition being requested, Cassandra employs a data structure known as a bloom filter.
- Bloom filters are a probabilistic data structure that allows Cassandra to determine one of two possible states
 - The data definitely does not exist in the given file
 - The data probably exists in the given file.



Architecture

- Gossip protocol
 - Each node learn about cluster topology.
 - Communicate among nodes.
 - Detection of faulty nodes.
- Snitch
 - Snitch helps map IPs to racks and data centers.
 - This info is used for replica location and other tasks.

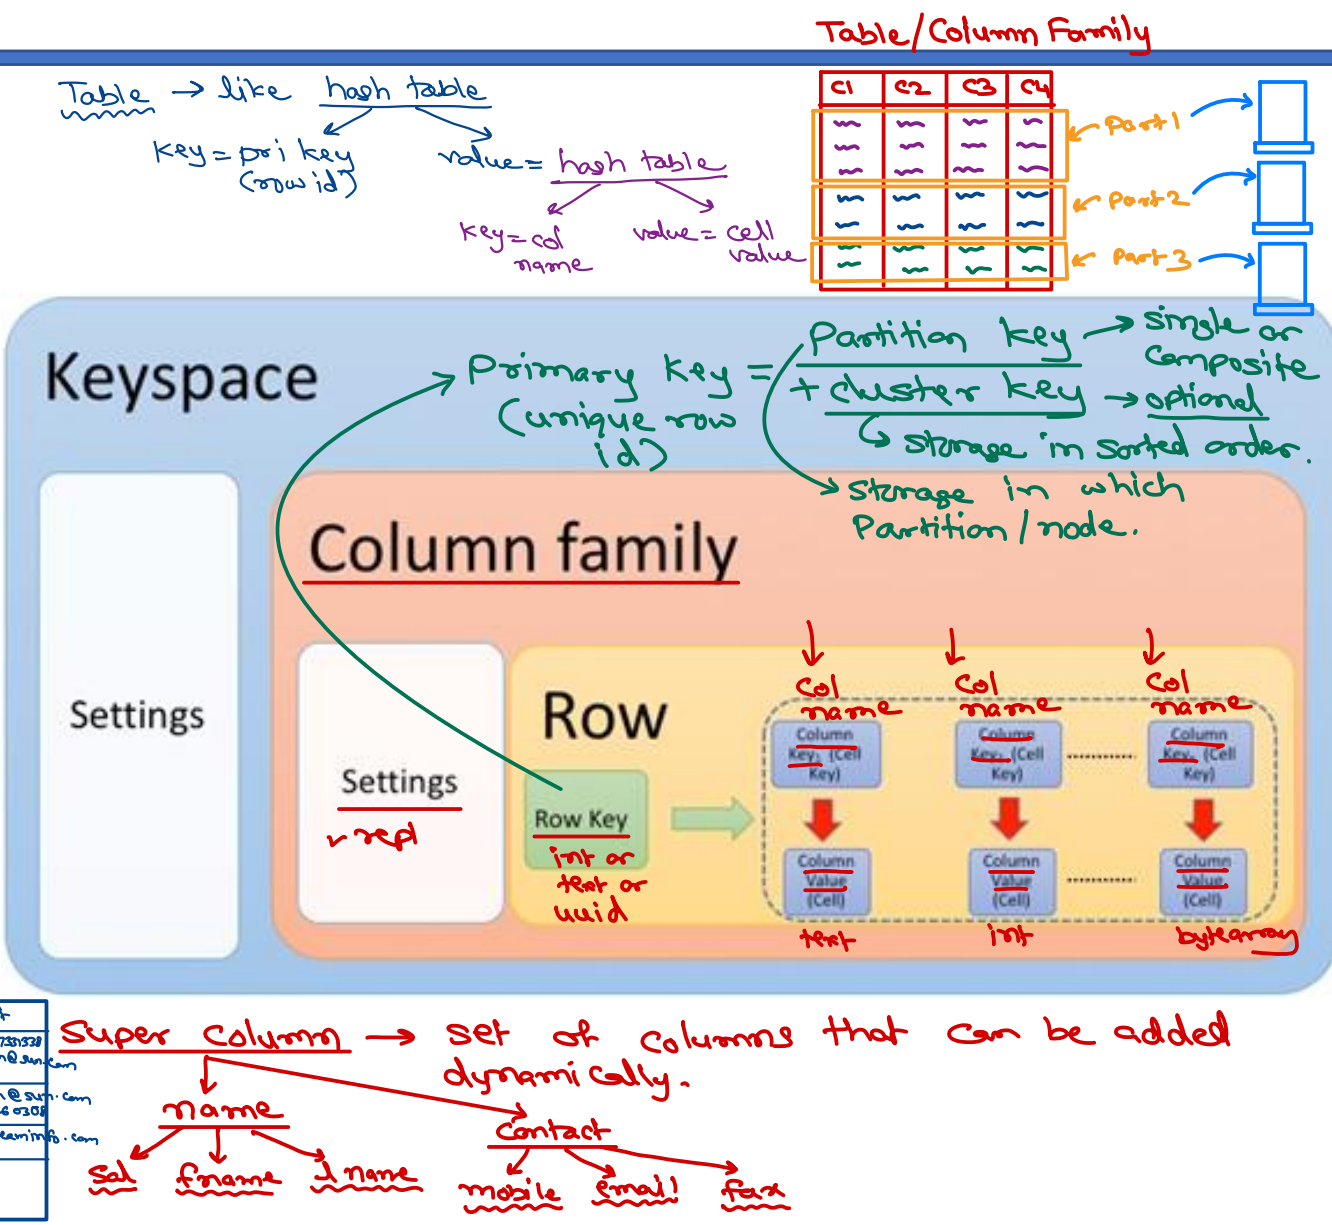


Data Model

- Basic data model is Rows and Columns and distributed across the nodes. → Partitions
 - Rows are distributed across nodes by Sharding on primary key. (Partition key)
 - Columns are distributed across nodes column groups.
- Each row is identified by unique key (primary key).
- One or more column families are contained inside Key-space (Table in CQL 3.0). → db schema
- A column family contain super-columns or columns.


people

name	contact
Sal: mr. Fname: nilesh Iname: ghule	mobile: 9827381928 email: nilesh@sunbeam.com
Fname: nikhil Iname: kudale	email: nikhil@sunbeam.com Fax: 2426 0302
Cname: Sunbeam Infotech	Site: Sunbeaminfo.com



Cassandra Keyspace

- DESCRIBE KEYSPACES;
- CREATE KEYSPACE dbda WITH replication = {'class':'SimpleStrategy', 'replication_factor' : 3};
- USE dbda;
- DESCRIBE TABLES;



- Strategy
 - Simple Strategy: Simple strategy is used in the case of one data center. In this strategy, the first replica is placed on the selected node and the remaining nodes are placed in clockwise direction in the ring without considering rack or node location.
 - Network Topology Strategy: This strategy is used in the case of more than one data centers. In this strategy, you have to provide replication factor for each data center separately.



CQL Data types

- ascii - US-ascii character string
- bigint - 64-bit signed long ints
- blob - Arbitrary bytes in hexadecimal
- boolean - True or False
- counter - Distributed counter values
64 bit
- decimal - Variable precision decimal
- double - 64-bit floating point
- float - 32-bit floating point
- frozen - Tuples, collections, UDT
containing CQL types
- inet - IP address in ipv4 or ipv6 string
format
- int - 32 bit signed integer
- list - Collection of elements
- map - JSON style collection of elements
- set - Sorted collection of elements
- text - UTF-8 encoded strings
- timestamp - ID generated with
date+time as int/string
- timeuuid - Type 1 uuid
- tuple - A group of 2,3 fields
- uuid - Standard uuid (128-bit)
- varchar - UTF-8 encoded string
- varint - Arbitrary precision integer



RDBMS

vs

Cassandra

- | | |
|-----------------------|--------------------------|
| ✓ Database/Schema | ✓ keyspace |
| ✓ Table | ✓ Table or Column Family |
| ✓ Row | ✓ Row |
| ✓ Primary Key | ✓ Row Key or Primary Key |
| ✓ Column Name | ✓ Column Name |
| ✓ Column value | ✓ Column or Cell value |
| ✓ Foreign keys, joins | ✓ Collections |
| ✓ Indexes | ✓ Indexes |





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

